

RoboCall AI

Phase 1 build status and carrier recommendation

14 September 2026 · Application: **built, tested, running** · Calling: **blocked pending a funded carrier account**

Summary. The dashboard and calling engine are complete and pass 203 automated checks. Google Voice cannot place these calls — it has no API and, more decisively, no audio path an AI could speak through. We built against Twilio instead, and the integration is correct: calls reach Twilio and are rejected there, because a trial account will not speak a script you wrote. The remaining decision is which carrier to fund — **Twilio at roughly \$20 with no code changes, or Telnyx at about half the per-minute cost.**

1 How the application works

A FastAPI server hosts both a JSON API and a plain HTML/CSS/JavaScript dashboard — no build step, no framework. It runs with one command and stores data in a local SQLite database. A campaign firing does not place N calls; it writes N rows to a task table that a background loop drains. That separation is what makes pacing, quiet hours and late opt-outs possible.

Stage	What happens
Compose	Write the message, or let the AI draft it from a one-line brief. A live preview shows what the recipient hears, including the auto-appended opt-out line.
Load	Paste numbers, import CSV files , or drag one onto the field. Any format normalises to E.164. A counter reports callable numbers, duplicates folded and bad lines before you submit.
Schedule	Once, hourly, daily or weekly, in a chosen timezone.
Dispatch	A daemon thread dials at 12 calls/minute — bursts are how numbers get flagged as spam. Runs whether or not a browser is open.
Reconcile	Outcomes polled back into the call log with the carrier's call ID.

Compliance is enforced in code, because the FCC's February 2024 ruling places AI-generated voices squarely under the TCPA: a consent flag per contact (unconsented numbers queue but never dial), caller identification and an opt-out line on every script, a per-recipient calling window derived from area code, a do-not-call list checked at both queue time and dial time, rate limiting, and a global kill switch. **203 automated checks** cover this against a fake carrier — 74 on the happy path, 129 adversarial. All pass.

2 Why Google Voice is not an option

The original brief proposed giving an AI agent access to a Google Voice number. This was investigated first, and it cannot work for three independent reasons.

- **There is no API.** Google retired it in 2014. What carries that name today is a Workspace *administrative* API that provisions licences for employees — no endpoint originates a call. The unofficial libraries are broken; Google changed its login flow and no working replacement exists.

- **It is prohibited.** Google Voice's Acceptable Use Policy forbids automated and bulk calling outright. The realistic outcome is a banned number partway through a campaign.
- **It is architecturally incapable — the decisive reason.** Click-to-call rings **your** phone first, then bridges you to the other party. Even with flawless automation there is no audio path for a synthesised voice to speak through, and because it occupies your physical handset it can place exactly one call at a time. The goal is mass announcement to a customer list. This cannot do it at any scale.

The number is still useful. It is already wired in as the **opt-out callback line** in every script — it takes the call, records a voicemail, and leaves the documented trail of removal requests the TCPA expects. It simply cannot be the thing that places the calls.

3 What we tried: Twilio

Twilio is reached in **direct mode** — the whole script rides along with the call-creation request as inline TwiML, which is why the prototype runs on a laptop with no tunnel or deployment. The carrier sits behind a single module, so the rest of the system is carrier-agnostic. Every row below is a genuine attempt against the production Twilio API using real credentials.

WHEN	CAMPAIGN	NUMBER	STATUS	LENGTH	TWILIO SID	DETAIL
Sep 13, 11:35 AM	Test call	(425) 285-5888	suppressed	—	—	
Sep 13, 11:34 AM	Test call	(425) 285-5888	campaign deleted	—	—	
Sep 13, 11:34 AM	Test call	(425) 285-5888	campaign deleted	—	—	
Sep 13, 11:34 AM	Test call	(425) 285-5888	campaign deleted	—	—	
Sep 13, 11:33 AM	Test call	(425) 285-5888	Failed	—	—	Twilio trial accounts do not accept the inline TwiML that direct mode sends, so no call can be placed from a trial in this mode. Either upgrade the account (Console -> Billing, add funds), or set a Public base URL in Settings to switch to webhook mode, which sends Twilio a URL instead. Twilio's own wording for this is "[0] Unable to create record: Invalid or disallowed parameters provided - trial accounts have limited parameter access, upgrade your account to unlock full functionality".
Sep 13, 11:31 AM	Test call	(425) 285-5888	Failed	—	—	[0] Unable to create record: Invalid or disallowed parameters provided - trial accounts have limited parameter access, upgrade your account to unlock full functionality

Figure 1 — The call log after live testing. The two **Failed** rows are Twilio rejecting the call; the upper one shows the same failure after the app was taught to explain it. Campaign deleted rows are queued calls cancelled with their test campaign, and suppressed is the do-not-call list working — the number was on it, so the engine refused to dial. The **Twilio SID** column is empty throughout: no call was ever created to have an ID.

Twilio's response to the script:

```
[0] Invalid or disallowed parameters provided - trial
accounts have limited parameter access, upgrade your account
```

Twilio does not say which parameter. It was isolated by changing one variable at a time: the inline script was rejected with answering-machine detection and again without it, while Twilio's own hosted demo URL was accepted. Per Twilio's documentation a trial call may set only the destination, a status callback, and one of four Twilio-hosted templates. **A trial will dial a phone, but it will not say words you wrote** — which is the entire product. The account reads Trial, \$0.00 balance, 0 numbers owned, 0 verified recipients.

4 The roadblock, and the choice

This is not a defect in the software and not a setting we have missed. Every major carrier gates custom automated speech behind a funded account, because that is precisely the capability fraudulent robocalling abuses. An account has to be funded; there are two sensible ways to do it.

	Twilio	Telnyx — recommended
To unblock	Add ~\$20, buy a number	Fund the account, order a number
Outbound voice, US	\$0.0140 / min	~\$0.0070 / min — about half
Local number	\$1.15 / month	\$1.00 upfront + \$1.00 / month
10,000 calls @ 30s	~\$70	~\$35
100,000 min / month	~\$1,400	~\$700 — \$8,400/year saved
Engineering work	None — already built and tested	A second provider behind the existing interface

Telnyx is the recommendation. It permits machine-generated speech with your own text — the thing Twilio's trial forbids — subject to a spoken prefix identifying the call as automated, which for a compliance-sensitive product is close to the caller identification the TCPA wants anyway. And it costs roughly half as much per minute: immaterial at prototype volume, the dominant line item as campaigns scale. The migration is smaller than it sounds — Telnyx speaks TeXML, compatible with the TwiML this application already generates, so the script builder, dispatcher, compliance engine, CSV import, scheduling and all 203 tests stay exactly as they are. An API key is already created and confirmed working against the account.

One caveat, stated plainly. Telnyx's documentation promises trial accounts \$5 in testing credit. That credit did **not** appear on the account created for this project — it reads \$0.00. Telnyx is recommended on per-minute cost and its willingness to speak custom text, not on a free tier. Twilio remains the faster option if zero engineering time matters more than the per-minute rate.

5 Next steps

#	Action	Owner
1	Fund the chosen carrier and order one voice-capable number.	You
2	Verify the test handset as a permitted recipient.	You
3	If Telnyx: build the provider behind the existing carrier interface.	Engineering
4	Place a live test call and confirm the script is spoken correctly.	Engineering

Before any real campaign: add authentication to the dashboard — every route is open today, fine on a laptop and unsafe on a public URL — subscribe to the National Do Not Call Registry and scrub against it, and retain signed consent records. Washington, where the test handset lives, is among the stricter state jurisdictions.

Shortly after: webhook mode with a public URL unlocks press-9-to-opt-out and hanging up on voicemail.